



BAHIR DAR INSTITUTE OF TECHNOLOGY

DEPARTMENT OF SOFTWARE ENGINEERING

Individual assignment

Cover Page

Installation of BlackArch operating system in vmware workstation,

Prepared by: Chalachew Akilew

ID:BDU1601197

Section :B

Course: Operating system and system programming

Table of Contents (TOC) for OSSP Individual Assignment

1:Introduction

- 1.1 Background of the Project
- 1.2 Motivation

2:Objectives of the Assignment

3:Historical Background of the Operating System (BlackArch Linux)

- 3.1 Evolution from Arch Linux
- 3.2 Purpose and Development Timeline

4:System Requirements

- 4.1 Hardware Requirements
- 4.2 Software Requirements

Installation of BlackArch Linux in VMware Workstation

- 5.1 Pre-Installation Setup
- 5.2 Step-by-Step Installation Process
- 5.3 Screenshots of the Installation
- 5.4 User Account Setup (Full Name Used)

Issues Faced During Installation

- 6.1 Network/Internet Problems
- 6.2 Package/Dependency Errors

7:Solutions to Installation Problems

Filesystem Support in BlackArch Linux

- 8.1 Supported Filesystems (ext4, Btrfs, etc.)
- 8.2 Recommended Filesystem for Installation
- 8.3 Why ext4 Was Chosen

Advantages and Disadvantages of BlackArch Linux

- 9.1 Advantages
- 9.2 Disadvantages

Virtualization in Modern Operating Systems

- 10.1 What is Virtualization?
- 10.2 Why Use Virtualization?
- 10.3 How Virtualization Works
- 10.4 Hypervisors and Types
- 10.5 Benefits in OS Development

Implementation of the `getuid()` System Call

- 11.1 Objective
- 11.2 Editing the Kernel Source
- 11.3 System Call Code Snippet
- 11.4 Compilation and Insertion of Module
- 11.5 Test Program Code
- 11.6 Screenshots and Terminal Output

Issues and Challenges Faced During System Call Implementation

- 12.1 Compilation Errors
- 12.2 Syscall Number Conflicts

Solutions and Fixes Applied

Evaluation Criteria Reflections

- 14.1 Installation Clarity
- 14.2 Historical Understanding
- 14.3 Virtualization Concepts
- 14.4 System Programming Experience
- 14.5 Filesystem Knowledge

Conclusion

- 15.1 Lessons Learned
- 15.2 Summary of Work Done

Future Outlook and Recommendations

- 16.1 Suggestions for New Learners
- 16.2 Plans for Further System Exploration

Appendices

- 18.1 Installation Screenshots
- 18.2 Code Listings
- 18.3 Terminal Logs

1:Introduction

1.1 Background of the Project

In the rapidly evolving field of software engineering, understanding how operating systems function at both user and kernel levels is essential. Operating System and System Programming (OSSP) is a foundational course that bridges the gap between software development and system-level engineering by focusing on low-level system behavior, OS architecture, and kernel-level programming.

This project was designed to provide hands-on experience with advanced system concepts through three core tasks: installing a specialized Linux distribution (BlackArch Linux) on a virtualized environment (VMware Workstation), analyzing the role of virtualization in modern computing, and implementing a custom system call within the Linux kernel. Each task is tailored to simulate real-world system-level challenges encountered by OS developers and security professionals.

BlackArch Linux was chosen because it is a powerful, Arch-based distribution tailored for security research. It presents a unique learning curve and environment ideal for exploring file systems, installation issues, and customization. Virtualization using VMware allows for isolated testing without affecting the host system, making it

perfect for experimentation. Implementing a system call further enhances understanding of how OS kernels interact with applications at the lowest level.

This project not only strengthens technical knowledge but also fosters deeper insight into systems programming, security, and performance optimization—key areas in modern operating system design and usage.

1.2 Motivation

The motivation behind this project stems from the need to gain practical experience in real-world system-level tasks that go beyond theoretical classroom knowledge. As a second-year Software Engineering student, it's essential to not only understand how operating systems work conceptually but also to interact with them directly through installation, customization, and programming at the kernel level.

This project offers the opportunity to explore key components of operating systems — from virtualization and filesystem structures to kernel-level programming through system call implementation. Installing and working with BlackArch Linux, a security-focused Arch-based distribution, introduces challenges that enhance problem-solving, troubleshooting, and Linux administration skills.

Furthermore, virtualization is a fundamental part of today's computing landscape, especially in cloud infrastructure and development environments. Understanding how virtualization works and being able to implement and manage it is an essential skill for any aspiring system engineer or developer.

Implementing a custom system call like `getuid()` helps bridge the gap between user space and kernel space programming — a rare and valuable learning experience. It demonstrates how high-level applications interact with low-level OS services, fostering a better understanding of system calls, memory management, and kernel internals.

This project is not only an academic requirement but also a stepping stone toward mastering operating systems and system-level programming, which are crucial for careers in systems development, cybersecurity, and DevOps.

2. Objectives of the Assignment

The primary objective of this OSSP (Operating Systems and System Programming) assignment is to develop foundational and practical knowledge of modern operating systems through hands-on experience in installation, virtualization, and system-level programming. This includes:

To install and configure a specialized Linux operating system (BlackArch) in a virtual environment using VMware Workstation, and to document the process effectively with real issues and solutions.

To understand and analyze virtualization technology in modern OS environments — specifically, the what, why, and how of virtualization, hypervisors, and their impact on development and system efficiency.

To implement a custom system call (getuid) within the Linux kernel, including editing kernel source files, compiling the kernel module, and testing the call via user-level applications.

To investigate the supported filesystems in BlackArch Linux, understand their usage, and justify the most appropriate choice for installation.

To enhance problem-solving skills by identifying and resolving system-level issues encountered during installation and system programming tasks.

To bridge theoretical knowledge with real-world applications, encouraging deeper exploration of system architecture, security, and kernel interaction.

3:Historical Background of the Operating System (BlackArch Linux)

3.1 Evolution from Arch Linux

BlackArch Linux is a direct descendant of **Arch Linux**, a lightweight, flexible, and minimalist Linux distribution known for its simplicity and full user control. Arch Linux follows a **rolling release model**, meaning users always have access to the latest stable packages without needing to perform major upgrades.

Recognizing the stability and modularity of Arch Linux, a group of security researchers and developers created BlackArch as an **extension of Arch**, specifically tailored for **penetration testing, ethical hacking, and security auditing**. Instead of building a new OS from scratch, they took advantage of Arch's powerful foundation, package manager (pacman), and community infrastructure.

Initially, BlackArch started as an **unofficial repository** of security tools that could be added to an existing Arch installation. Over time, due to growing popularity and demand for a ready-to-use system, the developers released a **custom ISO** that included a preconfigured desktop environment, window manager, and hundreds of security tools — all while preserving Arch's lightweight and DIY philosophy.

Today, BlackArch has evolved into a full-fledged **penetration testing distribution** with over **2,800 tools** categorized by function, while still remaining compatible with regular Arch Linux systems. It provides users the flexibility to either install it as a standalone OS or use it as an overlay on existing Arch installations — demonstrating its strong and ongoing connection to its parent distribution.

3.2 Purpose and Development Timeline

1:Purpose of BlackArch Linux:

BlackArch Linux is a penetration testing and security-focused distribution based on Arch Linux. Its primary purpose is to provide a comprehensive set of tools for security researchers, penetration testers, and system administrators.

2:Development Timeline:

Initial Development:

BlackArch Linux was first introduced in 2013 by a group of security researchers. It was developed as a way to provide a highly customizable and lightweight alternative to other penetration testing distributions like Kali Linux.

Milestones in Development:

2013: The first release of BlackArch was made available.

2014: BlackArch expanded its repository and included more security tools.

2015: Integration of the BlackArch tools into the Arch Linux ecosystem and regular updates to the toolset.

Present Day: BlackArch is continually updated with new tools, contributing to its reputation as one of the leading choices for cybersecurity professionals.

Rolling Release Model:

BlackArch follows a rolling release model, meaning it is continuously updated with the latest versions of tools and software. This approach ensures users always have access to the most recent developments in security tools.

Community and Contributions:

The development of BlackArch Linux has always been supported by a growing community of developers and contributors. You could mention how open-source contributions have played a vital role in the growth and success of the distro.

Conclusion:BlackArch Linux has evolved from a small project into one of the most widely recognized and used penetration testing distributions. Its continuous development ensures it stays relevant to the needs of security professionals.

4:System Requirements

i. Hardware:

- Intel Core i5 or higher
- Minimum 8GB RAM
- 40GB free disk space

ii. Software:

- VMware Workstation 17 Pro
- BlackArch Linux ISO
- Internet connection for package downloads

Here is a detailed explanation for **Section 5: Installation of BlackArch Linux in VMware Workstation**, specifically focusing on **Section 5.1: Pre-Installation Setup**. This will guide you through the necessary steps and preparations before beginning the installation of BlackArch Linux in VMware Workstation, tailored to your assignment's requirements.

5: Installation of BlackArch Linux in VMware Workstation

5.1 Pre-Installation Setup

Before starting the installation process of **BlackArch Linux** on **VMware Workstation**, there are a few **pre-installation steps** we need to follow to ensure that our system is prepared, and the virtual machine (VM) is correctly configured.

Here's a step-by-step guide for preparing the pre-installation setup:

1. Download the BlackArch ISO File

Objective: Obtain the official BlackArch Linux installation image.

How to do it:

Visit the [official BlackArch website](#).

Select the appropriate **64-bit ISO image** (since BlackArch is only available for 64-bit systems).

Download the ISO file to your local machine. Make sure we download the latest version to avoid using outdated tools.

2. Install VMware Workstation (or VMware Player)

Objective: Set up the virtualization software that will run BlackArch Linux in a virtualized environment.

How to do it:

If we haven't already installed **VMware Workstation**, you can download it from the [VMware website](#).

Follow the installation instructions provided by VMware to install the software on your system.

VMware Workstation Pro is recommended for advanced features, but **VMware Player** can also be used for non-commercial purposes.

System Requirements:

Ensure that our host machine has the minimum required resources (RAM, storage, and CPU) to run VMware Workstation smoothly.

3. Create a New Virtual Machine

Objective: Set up a new virtual machine in VMware Workstation for installing BlackArch Linux.

How to do it:

Launch **VMware Workstation** and click on "**Create a New Virtual Machine**".

Select "**Typical (recommended)**" for a simplified setup.

Choose "**Installer disc image file (iso)**" as the installation media, and then browse to the location of the **BlackArch ISO** you downloaded.

Click **Next** to continue.

4. Configure Virtual Machine Settings

Objective: Configure hardware settings for the virtual machine to optimize the installation and performance of BlackArch Linux.

How to do it:

Memory (RAM): Allocate at least **2 GB** of RAM to the VM. This is recommended for smoother operation during the installation and usage of security tools.

Processor (CPU): Allocate at least **2 core** to the virtual machine (2 cores are preferred for better performance).

Disk Space: Choose to create a new **virtual disk** with a size of at least **20 GB** to accommodate the BlackArch installation and additional tools.

Disk Type: Choose **SCSI** as the disk type for better performance and compatibility with VMware.

Ensure that the **Network Adapter** is set to **NAT** or **Bridged** for internet access.

5. Adjust Virtual Machine Options (Optional)

Objective: Fine-tune VM settings for better performance or customization.

How to do it:

Enable Virtualization Extensions: In the VMware settings, ensure that **virtualization extensions** (Intel VT-x or AMD-V) are enabled in the VM's processor settings if our host machine supports it.

Enable 3D Graphics: If we plan to use a graphical user interface (GUI), enable **3D graphics** for better graphical performance.

6. Prepare for Network Connectivity

Objective: Ensure the virtual machine will have proper internet access during installation.

How to do it:

Set the **Network Adapter** to either **Bridged** or **NAT** to allow the virtual machine to connect to the internet.

If using **Bridged**, the VM will be treated as a separate device on the local network, making it easier to install packages from the BlackArch repository.

If using **NAT**, the VM will share the host machine's IP address and be able to connect to the internet.

7. Check VMware Compatibility and Ensure Proper Settings

Objective: Verify that the virtual machine settings align with the recommended specifications for BlackArch Linux.

How to do it:

Double-check the **CPU**, **RAM**, **Disk Space**, and **Network Adapter** settings to ensure they meet or exceed the requirements for installing BlackArch.

Ensure that the **BlackArch ISO** file is properly mounted to the virtual CD/DVD drive.

Conclusion:

The **pre-installation setup** involves downloading the BlackArch Linux ISO, configuring a new virtual machine with the appropriate resources, ensuring network connectivity, and fine-tuning VM settings for optimal performance. These steps are essential in preparing VMware Workstation to run BlackArch Linux smoothly.

Once these preparations are complete, we will be ready to proceed with the actual **installation** of BlackArch Linux, as outlined in **Section 5.2**.

This detailed guide for **Pre-Installation Setup** will ensure that our environment is correctly prepared for the **BlackArch Linux installation** in VMware Workstation.

Here is a clear and well-organized explanation for **Section 5.2: Step-by-Step Installation Process** of BlackArch Linux in VMware Workstation, crafted to meet OSSP assignment requirements and criteria.

5.2 Step-by-Step Installation Process

Once the **pre-installation setup** is complete and the BlackArch ISO is mounted in VMware Workstation, follow these steps to **install BlackArch Linux** inside the virtual machine:

Step 1: Boot the Virtual Machine

Start the virtual machine in VMware Workstation.

The BlackArch boot menu will appear.

Choose BlackArch Linux (x86_64) or Boot BlackArch and press **Enter**.

The system will boot into a **live environment** (command line interface).

Step 2: Set Keyboard Layout (Optional)

By default, the keyboard layout is set to US.

If you use a different layout, you can set it using:

```
loadkeys <layout>
```

Example: `loadkeys uk`

Step 3: Connect to the Internet

Verify the network connection with

```
ping -c 3 google.com
```

If no internet:

Run `ip link` to check interface names (usually `eth0` or `enp0s3`).

Use DHCP to obtain an IP:

```
dhclient <interface_name>
```

Step 4: Sync System Clock

Synchronize the system time with:

```
timedatectl set-ntp true
```

Step 5: Partition the Disk

List available disks:

```
fdisk -l
```

Use `cfdisk` or `fdisk` to create partitions:

```
cfdisk /dev/sda
```

Recommended partitions:

`/boot` – 512MB

`swap` – 1GB or more

`/` – Remaining space (root)

Set filesystem type (ext4 recommended for root and boot).

Step 6: Format Partitions

Format created partitions:

```
mkfs.ext4 /dev/sda1 # boot
mkfs.ext4 /dev/sda3 # root
mkswap /dev/sda2
swapon /dev/sda2
```

Step 7: Mount Partitions

Mount the root and boot partitions:

```
mount /dev/sda3 /mnt
mkdir /mnt/boot
mount /dev/sda1 /mnt/boot
```

Step 8: Install the Base System

Install core BlackArch Linux packages:

```
pacstrap /mnt base base-devel linux linux-firmware vim nano
```

Step 9: Generate fstab

Create a file system table to auto-mount drives:

```
genfstab -U /mnt >> /mnt/etc/fstab
```

Step 10: Chroot into the New System

Change root into the new installed system:

```
arch-chroot /mnt
```

Step 11: Set Timezone and Locale

Set the timezone:

```
ln -sf /usr/share/zoneinfo/Region/City /etc/localtime
```

```
hwclock --systohc
```

Example: `ln -sf /usr/share/zoneinfo/Africa/Addis_Ababa /etc/localtime`

Configure locale:

```
nano /etc/locale.gen
```

Uncomment: `en_US.UTF-8 UTF-8` Then run:

```
locale-gen  
echo "LANG=en_US.UTF-8" > /etc/locale.conf
```

Step 12: Set Hostname

Assign a hostname:

```
echo "blackarch-vm" > /etc/hostname
```

Step 13: Set Root Password

Create a root password:

```
passwd
```

Step 14: Install and Configure Bootloader

Install GRUB:

```
pacman -S grub  
grub-install --target=i386-pc /dev/sda  
grub-mkconfig -o /boot/grub/grub.cfg
```

Step 15: Exit, Unmount, and Reboot

Exit from chroot:

```
exit
```

Unmount all partitions:

```
umount -R /mnt
```

Reboot the system:

```
reboot
```

Step 16: First Boot

Once rebooted, the system will start from the newly installed BlackArch Linux.

Login as root with the password you set.

Optional: Install BlackArch Tools

After booting into your installed system, you can install BlackArch tools using:

```
pacman -Syyu  
pacman -S blackarch
```

Summary:

This step-by-step process ensures a successful and smooth installation of BlackArch Linux in VMware Workstation. It covers disk partitioning, base system installation, locale configuration, and bootloader setup—tailored for virtualization. Following this process ensures your system is correctly prepared for further experimentation and system programming tasks, like adding system calls.

Here's a well-written explanation for **Section 5.4: User Account Setup (Full Name Used)** tailored for your OSSP documentation and formatted according to assignment expectations:

5.4 User Account Setup (Full Name Used)

After successfully installing BlackArch Linux and rebooting into the system, it is essential to create a **non-root user** account for better security and usability. The following steps outline how the user account was set up using my full name:

Step 1: Log in as Root

Upon first boot, log in using the root account with the password set during installation.

```
login: root  
Password: *****
```

Step 2: Create a New User (Using Full Name)

To create a new user account with your full name, use the following command:

```
useradd -m -G wheel -s /bin/bash chalachewakilew
```

-m: Creates a home directory for the user.

-G wheel: Adds the user to the *wheel* group (for sudo permissions).

-s /bin/bash: Sets the default shell to Bash.

chalachewakilew:

Username based on my **full name: Chalachew Akilew**

Step 3: Set User Password

Set a password for the newly created user:

```
passwd chalachewakilew
```

I will be prompted to enter and confirm the password.

Step 4: Enable Sudo for Wheel Group

To allow the user to perform administrative tasks using `sudo`, edit the **sudoers** file:

```
EDITOR=nano visudo
```

Find the following line and **uncomment it** by removing the #:

```
# %wheel ALL=(ALL:ALL) ALL
```

After editing, save and exit (Ctrl + X, then Y and Enter in nano).

Step 5: Switch to New User

I can now switch to the newly created user account:

```
su - chalachewakilew
```

```
error: failed to commit transaction (invalid or corrupted package (PGP signature))
Errors occurred, no packages were upgraded.
Installing for i386-pc platform.
grub-install: error: failed to get canonical path of 'airoutfs'.
/usr/bin/grub-probe: error: failed to get canonical path of 'airoutfs'.
[ blackarch ~/Desktop ]# useradd -m -G wheel -s /bin/bash chalachew_akilew
[ blackarch ~/Desktop ]# passwd chalachew_akilew
passwd chalachew
passwd: user 'chalachew' does not exist
[ blackarch ~/Desktop ]# passwd chalachew_akilew
New password:
Retype new password:
Sorry, passwords do not match.
passwd: Failed preliminary check by password service
passwd: password unchanged
[ blackarch ~/Desktop ]# passwd chalachew_akilew
New password:
Retype new password:
passwd: password updated successfully
[ blackarch ~/Desktop ]# whoami
root
[ blackarch ~/Desktop ]# su - chalachew_akilew
[ blackarch ~ ]# whoami
chalachew_akilew
[ blackarch ~ ]# sudo pacman -S linux-source

We trust you have received the usual lecture from the local System
Administrator. It usually boils down to these three things:

#1) Respect the privacy of others.
#2) Think before you type.
#3) With great power comes great responsibility.

For security reasons, the password you type will not be visible.

Terminal - chalachew_akilew@blackarch:~
```

Or reboot and log in directly using:

login: chalachewakilew

Password: *****

Verification

To test whether the user has sudo privileges, run:

sudo pacman -Syu

```
Terminal - chalachew_akilew@blackarch:~
:: Synchronizing package databases...
core is up to date
extra is up to date
multilib is up to date
blackarch is up to date
resolving dependencies...
looking for conflicting packages...

Packages (1) archlinux-keyring-20250123-1
Total Download Size: 1.18 MiB
Total Installed Size: 1.68 MiB
Net Upgrade Size: 0.07 MiB

:: Proceed with installation? [Y/n] Y
:: Retrieving packages...
archlinux-keyring-20250123-1-any 1206.9 KiB 583 KiB/s 00:02 [#####] 100%
(1/1) checking keys in keyring [#####] 100%
(1/1) checking package integrity [#####] 100%
(1/1) loading package files [#####] 100%
(1/1) checking for file conflicts [#####] 100%
:: Processing package changes...
(1/1) upgrading archlinux-keyring [#####] 100%
==> Appending keys from archlinux.gpg...
==> Locally signing trusted keys in keyring...
-> Locally signed 1 keys.
==> Importing owner trust values...
gpg: inserting ownertrust of 4
==> Disabling revoked keys in keyring...
-> Disabled 11 keys.
==> Updating trust database...
gpg: marginals needed: 3 completes needed: 1 trust model: pgp
gpg: depth: 0 valid: 1 signed: 6 trust: 0-, 0q, 0n, 0m, 0f, 1u
gpg: depth: 1 valid: 6 signed: 101 trust: 1-, 0q, 0n, 5m, 0f, 0u
gpg: depth: 2 valid: 78 signed: 28 trust: 78-, 0q, 0n, 0m, 0f, 0u
gpg: next trustdb check due at 2025-06-16
:: Running post-transaction hooks...
(1/2) Reloading system manager configuration...
(2/2) Arming ConditionNeedsUpdate...
[ blackarch ~ ]#
```

If prompted for a password and the command executes successfully, the user setup is complete.

Summary

A secure and personalized user account has been created using my full name, **Chalachew Akilew**, with proper permissions for system management via the wheel group and sudo. This enhances both usability and security, following best practices in Linux system administration.



6:Issues Faced During Installation:

- BlackArch installer failed due to missing dependencies.
- Internet connection failed inside the virtual machine.

7:Solutions to Installation Problems

- Updated package manager using pacman -Syu before installation.
- Configured bridged networking for internet access.

8:Filesystem Support in BlackArch Linux

It seems like I can't do more advanced data analysis right now. Please try again later.

In the meantime, I can give you the exact content for **Section 8.1: Supported Filesystems** so you can copy and paste it directly into your document:

8.1 Supported Filesystems (ext4, Btrfs, etc.)

BlackArch Linux, being based on Arch Linux, supports a wide range of modern and legacy filesystems. These filesystems determine how data is stored, organized, and accessed on storage devices. Some of the key filesystems supported by BlackArch include:

✓ **ext4 (Fourth Extended Filesystem)**

Most widely used and the default filesystem in most Linux distributions, including BlackArch.

Offers high performance, reliability, and journaling.

Supports large file sizes and volumes.

✓ **Btrfs (B-tree Filesystem)**

Modern copy-on-write filesystem designed for scalability.

Supports snapshots, subvolumes, RAID, compression, and data integrity checks.

✓ **XFS**

High-performance journaling filesystem used in enterprise environments.

Excellent for handling large files and parallel I/O operations.

✓ **FAT32 / exFAT**

Commonly used for USB drives and cross-platform compatibility with Windows.

Supported by Linux, but not ideal for Linux OS installations

✓ **NTFS**

Native to Windows, supported in Linux via the ntfs-3g driver.

Suitable for dual-boot or external drives shared with Windows.

✓ **ZFS**

Advanced filesystem known for high integrity, snapshots, and scalability.

Requires additional modules; less commonly used in BlackArch.

✓ **Other Filesystems (ReiserFS, JFS, HFS+, etc.)**

Available via optional kernel modules.

Used mainly for legacy or cross-platform compatibility.

8.3 Why ext4 Was Chosen

Great! Here's a strong explanation for **8.3: Why ext4 Was Chosen**, which clearly justifies your choice from a technical and practical perspective:

8.3 Why ext4 Was Chosen

For the installation of BlackArch Linux, **ext4** was chosen as the filesystem due to its balance of **performance, stability, compatibility, and ease of use**. As a second-year software engineering student working in a virtualized environment, choosing the most reliable and least error-prone filesystem was essential for a smooth installation and experimentation experience.

Key Reasons for Choosing ext4:

Default in Arch and BlackArch Linux

ext4 is the default filesystem supported by the Arch Linux family, making it the most compatible and well-integrated choice during the installation of BlackArch.

Performance Efficiency

It offers excellent read and write speeds with minimal CPU usage, which is especially important when running a virtual machine where resources are shared.

Data Safety via Journaling

ext4 includes journaling, which records changes before applying them. This protects against data corruption in case of unexpected shutdowns, making it ideal for environments where reliability matters.

Low Complexity and High Support

Unlike advanced filesystems like Btrfs or ZFS that require additional configuration or kernel modules, ext4 works out-of-the-box and is supported by all major Linux utilities, bootloaders, and system tools.

Stability and Proven Reliability

With over a decade of use in production systems, ext4 is a mature and trusted filesystem with a proven record of robustness and consistency — especially suitable for learners and professionals alike.

In conclusion, ext4 was selected because it ensured a **smooth, stable, and secure** installation process while offering full compatibility with BlackArch Linux and VMware. It was the most practical and technically sound choice for this project.

9: Advantages and Disadvantages of BlackArch Linux

9.1 Advantages:

Sure! Here's a strong, focused explanation for **9.1 Advantages** of using BlackArch Linux — especially in the context of your OSSP project:

9.1 Advantages

BlackArch Linux offers several advantages that make it a powerful choice for advanced users, ethical hackers, and cybersecurity researchers. These benefits also align well with the learning goals of system programming and operating system exploration.

1. Extensive Tool Collection

BlackArch includes over **2,800 penetration testing and security tools**, pre-installed and categorized for convenience. This eliminates the need to manually search and configure tools for ethical hacking, digital forensics, or reverse engineering.

2. Based on Arch Linux

As an Arch-based distribution, BlackArch inherits the **rolling release model**, giving users access to the **latest packages and updates** without needing to perform full upgrades — ideal for continuous development and research environments.

3. Lightweight and Minimal

Unlike some bulky security distributions, BlackArch offers a minimal installation with optional window managers, giving users full control over system resources and customization.

4. High Customizability

Advanced users can fully customize their environment, toolsets, user interface, and services — making it perfect for research, experimentation, and tailored workflows.

5. Strong Community and Documentation

Although BlackArch is not as mainstream as Kali Linux, it benefits from the broader Arch Linux ecosystem. Users can rely on the **Arch Wiki**, package manager (pacman), and a supportive community for troubleshooting and guidance.

6. Ideal for Virtualization and Testing

BlackArch works well in virtual environments (e.g., VMware, VirtualBox), making it suitable for **safe penetration testing** and **OS experimentation** without affecting the host system.

Overall, BlackArch provides a powerful and modular Linux environment that is perfect for advanced system programming, security research, and real-world operating system experimentation.

Absolutely! Here's a clear and honest explanation for **9.2 Disadvantages** of using BlackArch Linux, especially from the perspective of students or new users:

9.2 Disadvantages

While BlackArch Linux is a powerful and highly customizable distribution, it also comes with certain limitations and challenges that may impact its usability—especially for beginners or those unfamiliar with Arch-based systems.

1. Steep Learning Curve

BlackArch is not designed for Linux beginners. It lacks a graphical installer and assumes that users are already comfortable with the **Linux command line**, **manual partitioning**, and system configuration.

2. Manual Installation Process

Unlike distributions such as Kali Linux or Ubuntu, BlackArch requires users to manually configure most of the installation steps. This can lead to mistakes or confusion if not done correctly, especially during partitioning, network setup, or bootloader configuration.

3. Resource Intensive with Full Toolset

When all tools are installed, BlackArch can consume significant storage and memory. In constrained virtual environments, this can slow down the system and affect performance unless properly optimized.

4. Limited Official Documentation

Although BlackArch benefits from the Arch Wiki, its own documentation is minimal. This can make it harder to find specific guidance for BlackArch-specific issues compared to more mainstream distributions like Kali Linux.

5. Potential for Instability

Because it follows a **rolling release model**, package updates are frequent. Without careful update management, this can sometimes lead to **system breakage** or incompatibility with tools.

6. Not Ideal for Daily Use

BlackArch is tailored for penetration testing and security tasks — not for general-purpose computing. It lacks common desktop applications by default, making it unsuitable for users seeking an everyday OS.

Conclusion: While BlackArch is excellent for advanced users and specific use cases like ethical hacking and system programming, it may be **too complex or unstable for new users** or those seeking a general-use Linux environment.

Question 2:

10:Virtualization in Modern Operating Systems

Of course! Here's a clear, concise, and technically accurate explanation for:

10.1 What is Virtualization?

Virtualization is the process of creating a **virtual (software-based) version** of a physical computing resource, such as an operating system, server, storage device, or network. It allows a single physical machine to run **multiple independent systems or applications** by abstracting the underlying hardware.

In operating systems, virtualization enables the creation of **virtual machines (VMs)** isolated environments that each behave like a separate physical computer. These VMs

are managed by a **hypervisor**, which allocates hardware resources such as CPU, memory, and disk space to each VM.

Key Concepts:

Each virtual machine can run its own **operating system and applications**, regardless of the host OS.

The virtualized environment is **isolated**, meaning changes in one VM do not affect others.

Virtualization improves **resource utilization, system flexibility, and test/development workflows**.

Example:

Running **BlackArch Linux inside VMware Workstation** is a perfect example of virtualization — it allows testing a separate OS without changing or affecting the host system (e.g., Windows or Ubuntu).

Absolutely! Here's a solid and well-structured explanation for:

10.2 Why Use Virtualization?

Virtualization is widely used in both academic and professional environments because it offers **numerous advantages** over traditional physical systems. It simplifies system management, enhances resource efficiency, and allows safer testing and experimentation — all of which are crucial in operating system and system programming work.

1. Resource Optimization

Virtualization allows a single physical machine to run **multiple operating systems** and applications, making **better use of CPU, RAM, and storage**. This eliminates the need for separate hardware for each system.

2. Safe Testing Environment

It provides an **isolated and controlled environment**, which is ideal for experimenting with different OS installations, system call implementations, kernel modifications, and even malware testing — all without risking the host system.

3. Cost-Effective

Instead of purchasing multiple physical computers, virtualization lets users simulate different OS environments on one machine, reducing hardware and energy costs.

4. Flexibility and Portability

Virtual machines can be **easily created, cloned, backed up, and moved** across systems or platforms, making them highly adaptable for both development and deployment.

5. Supports Multiple OSes Simultaneously

Developers can run **Linux, Windows, macOS**, or any combination side-by-side for **cross-platform development and compatibility testing**.

6. Snapshot and Rollback Features

Hypervisors like VMware allow users to **take snapshots** of a system's current state and **roll back** if something breaks during experimentation. This is incredibly useful for learning and debugging.

In Summary: Virtualization is an essential tool for software engineers, system programmers, and cybersecurity students because it provides a **safe, flexible, and efficient** platform for learning, development, and testing.

10.3 How Virtualization Works

Virtualization works by introducing a **software layer called a hypervisor** between the physical hardware and the operating systems. This hypervisor abstracts hardware resources (CPU, memory, disk, network) and dynamically allocates them to **multiple virtual machines (VMs)** running simultaneously on the same physical machine.

Key Components:

1. Hypervisor (Virtual Machine Monitor)

A hypervisor is the core of virtualization. It manages and monitors virtual machines by controlling hardware access and resource sharing. There are two main types:

Type 1 (Bare-metal Hypervisor):

Runs directly on hardware. Example: VMware ESXi, Microsoft Hyper-V.

➤ Used in servers and data centers.

Type 2 (Hosted Hypervisor):

Runs on top of an existing host operating system. Example: VMware Workstation, VirtualBox.

➤ Ideal for desktop users and student projects.

2. Virtual Machine (VM)

A virtual machine is a software-based emulation of a physical computer. Each VM has:

Its own virtual CPU, memory, disk, and network interface

Its own **guest OS** (e.g., BlackArch, Ubuntu, Windows)

Complete isolation from other VMs and the host system

3. Virtual Hardware Abstraction

The hypervisor creates **virtual hardware** for each VM, so the guest OS thinks it's running on real hardware. For example:

Virtual CPU = mapped from physical CPU cores

Virtual RAM = allocated from host memory

Virtual Disk = mapped from host storage (e.g., .vmdk files)

4. Resource Scheduling and Management

The hypervisor **dynamically schedules** access to hardware resources among the VMs based on demand, limits, and priorities. This allows multiple VMs to run efficiently on a single system.

Example in Practice:

When we install **BlackArch Linux** in **VMware Workstation**, the VMware hypervisor:

Creates a virtual hard disk for BlackArch

Allocates RAM and CPU for the VM

Provides a virtual display, USB controller, and internet adapter

Lets we run BlackArch as if it were installed on a separate computer

:10.4 Hypervisors and Types

A **hypervisor**, also known as a **Virtual Machine Monitor (VMM)**, is a software layer that enables virtualization by allowing multiple virtual machines (VMs) to run on a single physical host. It manages the distribution of hardware resources such as CPU, memory, and storage among these VMs and ensures that they operate independently and securely.

There are two main types of hypervisors:

1. Type 1 Hypervisor (Bare-Metal)

This type runs **directly on the host's hardware**, without relying on a host operating system.

Examples: VMware ESXi, Microsoft Hyper-V, Xen, KVM (Kernel-based Virtual Machine)

Use Cases: Data centers, enterprise servers, cloud computing infrastructure

Advantages:

Offers **better performance** and efficiency

More **secure and stable**

Designed for **high-availability environments**

2. Type 2 Hypervisor (Hosted)

This type runs **on top of an existing operating system** (like Windows or Linux) and uses that OS to manage the hardware.

Examples: VMware Workstation, Oracle VirtualBox, Parallels Desktop

Use Cases: Personal computers, development and testing environments

Advantages:

Easier to set up and use

Ideal for **learning, testing, and development**

Can be installed on any desktop OS

Choosing Between Type 1 and Type 2

For professional, large-scale production environments → **Type 1** is ideal.

For educational, experimental, or home use → **Type 2** is more convenient.

In My OSSP project, **VMware Workstation** (a Type 2 hypervisor) is used to install **BlackArch Linux** on My host machine, making it perfect for OS experimentation without the need for dual-booting or additional hardware.

10.5 Benefits in OS Development (Virtualization in OS Development)

Virtualization plays a significant role in modern operating system (OS) development. It provides several key benefits that enhance the development, testing, and maintenance of operating systems. Here are the main benefits of virtualization in OS development:

1. Isolation of Environments

Virtualization allows developers to create isolated environments for testing different versions of an OS, configurations, or software without impacting the host system. This isolation ensures that any errors or crashes in the virtualized environment do not affect the primary operating system, providing a safe space for experimentation.

2. Efficient Resource Management

Virtual machines (VMs) allow multiple OS instances to run simultaneously on a single physical machine, optimizing hardware usage. By allocating specific resources (CPU, memory, disk space) to each virtual machine, developers can test how the OS behaves under varying resource conditions and workloads, making it easier to manage and utilize system resources effectively.

3. Cross-Platform Testing

Virtualization enables developers to test their OS or applications across different hardware platforms and operating systems without requiring multiple physical machines. Developers can set up VMs for testing on different OS types (Windows, Linux, macOS, etc.), enabling cross-platform compatibility testing and debugging.

4. Rapid Prototyping and Debugging

Virtualization speeds up the prototyping phase of OS development. Developers can rapidly create new VM instances, test different configurations, and debug without worrying about the setup time and hardware constraints. If a configuration causes issues, the developer can easily reset the virtual environment to a known stable state.

5. Cost Efficiency

Using virtual machines reduces the need for physical hardware for testing, which can be expensive and space-consuming. Developers can run multiple virtual environments on a single machine, significantly reducing hardware costs and maintenance overhead.

6. Snapshot and Rollback

Virtualization platforms allow developers to take snapshots of a system's state at any point in time. This feature makes it easy to revert to a previous working state if a change causes problems or crashes. Snapshots are invaluable when experimenting with new configurations, system changes, or OS updates.

7. Parallel Testing of Multiple OS Versions

With virtualization, developers can run different OS versions in parallel on the same machine. This is especially useful in OS development when testing backward compatibility, evaluating new features, or comparing the performance of different OS versions under similar conditions.

8. Simplified Deployment and Distribution

Virtualized environments allow easy deployment and distribution of operating systems or applications. Developers can distribute virtual machine images (e.g., VM snapshots or disk images) to testers, collaborators, or clients, ensuring that everyone is testing the same environment, minimizing discrepancies caused by hardware differences.

9. Security Testing

Virtual machines offer a secure testing environment where developers can simulate potential security vulnerabilities and exploits without putting the host system at risk. This is critical for OS development, as it helps ensure that the operating system is secure before release.

10. Simulating Complex Scenarios

Developers can simulate complex network and multi-system setups in virtual environments, making it easier to test distributed systems or network-intensive applications. Virtual machines can be networked together, replicating large-scale system architectures for thorough testing.

Conclusion: Virtualization accelerates OS development by providing developers with powerful tools to isolate, test, debug, and deploy operating systems and applications efficiently. It reduces hardware costs, enhances testing capabilities, and supports cross-platform development, making it an essential tool in modern OS development processes.

Question 4:

11.Implementation of the `getuid()` System Call

11.1 Objective

The objective of this section is to implement the `getuid()` system call in BlackArch Linux, a security-focused derivative of Arch Linux. The `getuid()` system call is a part of the Linux kernel and is responsible for returning the user ID (UID) of the currently logged-in user. This user ID is used to determine user privileges and manage access control within the system.

The goal of this implementation is to:

- Understand how system calls interact with the Linux kernel.

- Modify the kernel source code to add the `getuid()` system call.

- Ensure the correct return of the UID of the calling process, providing insight into the kernel's process management system.

Compile the kernel with the newly added system call and verify its functionality with a test program.

By implementing this system call, the assignment aims to deepen the understanding of system programming, specifically how the kernel handles requests from user space and how new system calls can be integrated into the existing kernel framework. The final implementation will involve modifying kernel files, compiling the kernel, and testing the system call using a simple C program that calls `getuid()`.

11.2 Editing the Kernel Source (Implementation of the `getuid()` System Call)

In order to implement the `getuid()` system call in BlackArch Linux, you need to modify the kernel source code. This involves adding the system call's functionality, defining its entry point, and ensuring it is properly registered within the kernel's system call table. Below are the key steps to follow for editing the kernel source:

Steps for Editing the Kernel Source:

Obtain the Kernel Source Code:

First, ensure that we have the kernel source code for BlackArch Linux. we can download the source code for the specific version of the kernel we are working with.

If we are working with a custom kernel or a specific version of Arch Linux's kernel, we can find it on the Arch Linux Git repository or use `pacman` to install the kernel source.

```
sudo pacman -S linux
```

Navigate to the Kernel Source Directory:

Once you have the kernel source code, navigate to the directory where the kernel is located. For example:

```
cd /usr/src/linux
```

Define the `getuid()` System Call:

The `getuid()` system call returns the real user ID of the calling process. The implementation involves modifying the appropriate source file to define the system call.

The system call is typically implemented in the kernel's `kernel/sys.c` file, where various system calls are defined.

Code for `getuid()` System Call:

```
asmlinkage long sys_getuid(void)
```

```
{
    return current->uid;
}
```

`current->uid` returns the real user ID of the process calling the system call. The `current` pointer is a structure that represents the current task (process) in the Linux kernel.

Register the System Call:

After defining the system call, it must be added to the system call table, which maps system call numbers to function pointers in the kernel.

Find the system call table for your kernel. Typically, this is in the file `arch/x86/entry/syscalls/syscall_32.tbl` for 32-bit systems or `arch/x86/entry/syscalls/syscall_64.tbl` for 64-bit systems.

Add an entry for the `getuid()` system call by choosing a free system call number (e.g., 233) and mapping it to the `sys_getuid()` function.

Example:

```
233    common    sys_getuid
```

Update the Header Files:

You need to add the function prototype for `sys_getuid()` to a header file. This will make the system call available to the user space programs that will interact with it.

Typically, this is done in the `include/linux/syscalls.h` file, where other system call prototypes are declared

Example:

```
asmlinkage long sys_getuid(void);
```

Rebuild the Kernel:

After modifying the kernel source files, you will need to recompile the kernel to include the new system call.

Run the following commands to build the kernel:

```
make menuconfig # (Optional) Configure kernel options
make            # Compile the kernel
sudo make install # Install the new kernel
```

Reboot into the New Kernel:

Once the kernel has been compiled and installed, reboot the system into the new kernel that includes the `getuid()` system call.

```
sudo reboot
```

By completing these steps, we will have successfully edited the kernel source to implement the `getuid()` system call. The next step is to test the system call to ensure that it works correctly.

11.3 System Call Code Snippet (Implementation of the `getuid()` System Call)

Here is the code snippet for the implementation of the `getuid()` system call in the Linux kernel:

1. System Call Definition

First, define the system call function `sys_getuid()` in the appropriate kernel source file. This function will return the real user ID (uid) of the calling process.

```
#include <linux/kernel.h> // For kernel-related functions
#include <linux/sched.h> // For current task (process)

asmlinkage long sys_getuid(void)
{
    return current->uid; // Return the real user ID of the calling process
}
```

Explanation:

`asmlinkage` is a macro used to define the calling convention for system calls in the Linux kernel.

`current` is a pointer to the current task (process) in the kernel. It is a global variable that points to the `task_struct` of the currently running process.

`current->uid` gives the real user ID of the calling process.

2. Register the System Call

Next, register the `sys_getuid()` system call in the system call table so that the kernel knows how to invoke it when requested by a user-space application.

For a 64-bit system, you would modify the `arch/x86/entry/syscalls/syscall_64.tbl` file:

```
233    common    sys_getuid
```

This entry maps the system call number 233 to the `sys_getuid()` function.

3. Header File Update

To make the system call available to user-space applications, declare the function prototype in the `include/linux/syscalls.h` header file:

```
asmlinkage long sys_getuid(void);
```

This declaration allows the kernel to recognize the `sys_getuid()` function when referenced in other kernel files.

4. Test Program Code Snippet

To test the `getuid()` system call after implementing it, you can write a simple C program that calls the system call. This program will call `getuid()` and print the returned user ID.

```
#include <stdio.h>
#include <unistd.h> // For syscall function
#include <sys/syscall.h> // For syscall numbers

#define __NR_getuid 233 // Define the syscall number for getuid

int main()
{
    // Call the getuid system call and print the returned user ID
    long uid = syscall(__NR_getuid);
    printf("The user ID is: %ld\n", uid);

    return 0;
}
```

Explanation:

The program uses the `syscall()` function to invoke the system call by its number (`__NR_getuid`).

The result is printed, which should be the real user ID of the process.

This completes the code snippets for both the kernel implementation of the `getuid()` system call and a simple user-space program to test it.

```
GNU nano 2.2 getuid.c
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/init.h>
#include <linux/syscalls.h>
#include <linux/uaccess.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Your Name");
MODULE_DESCRIPTION("A simple module for the getuid syscall");

asmlinkage long (*original_getuid)(void);

asmlinkage long new_getuid(void) {
    printk(KERN_INFO "New getuid() syscall called\n");
    return original_getuid();
}

static int __init getuid_init(void) {
    printk(KERN_INFO "getuid module loaded\n");
    original_getuid = (void *)kallsyms_lookup_name("sys_getuid");
    if (original_getuid) {
        *((unsigned long *)original_getuid) = (unsigned long)new_getuid;
        printk(KERN_INFO "getuid syscall hooked successfully\n");
    } else {
        printk(KERN_ERR "Unable to locate sys_getuid\n");
    }
    return 0;
}

static void __exit getuid_exit(void) {
    if (original_getuid) {
        *((unsigned long *)original_getuid) = (unsigned long)original_getuid;
        printk(KERN_INFO "getuid syscall unhooked\n");
    }
    printk(KERN_INFO "getuid module unloaded\n");
}

__exit(getuid_exit);
```

11.4 Compilation and Insertion of Module (Implementation of the getuid() System Call)

After defining and registering the `getuid()` system call in the kernel source code, the next step is to compile the kernel, install the new kernel, and ensure that the changes take effect. This involves several stages: compiling the kernel, inserting any necessary modules, and rebooting into the new kernel.

Steps for Compilation and Insertion of Module:

1. Install Necessary Dependencies

Before compiling the kernel, ensure that you have all the required development tools and dependencies installed. These tools are essential for building and configuring the kernel:

```
sudo pacman -S base-devel bc kmod cpio
```

2. Configure the Kernel (Optional)

If you want to modify any kernel options, you can run the configuration utility. This is an optional step if you need to customize kernel features.

```
make menuconfig
```

However, for the `getuid()` system call, no special kernel configuration is required unless you're making other modifications.

3. Compile the Kernel

Once the kernel source is ready and changes have been made, you need to compile the kernel. This can take some time, depending on the system's resources.

Build the Kernel: To compile the kernel, run the following command from the root of the kernel source directory:

```
make -j$(nproc)
```

The `-j$(nproc)` flag tells the system to use all available CPU cores to speed up the compilation process.

Build Kernel Modules: You may also need to build kernel modules if you're adding any. To build the modules:

```
make modules
```

Install the Kernel: After the kernel has been successfully compiled, you need to install it. This installs the newly compiled kernel and its associated modules.

```
sudo make modules_install  
sudo make install
```

4. Update Bootloader (If Needed)

If the kernel installation modifies the bootloader configuration (e.g., adding a new kernel version), you might need to update the bootloader. For example, if you're using GRUB, you can run:

```
sudo grub-mkconfig -o /boot/grub/grub.cfg
```

This will regenerate the GRUB configuration and ensure that the new kernel is included in the boot options.

5. Reboot the System

After the kernel has been compiled, installed, and the bootloader has been updated (if needed), reboot the system to load the new kernel.

```
sudo reboot
```

6. Verify the System Call

After rebooting, you can verify that the `getuid()` system call is working correctly by running a test program. The test program will call the system call and display the real user ID of the calling process.

Use the program created in Section 11.3 (Test Program Code Snippet) to call `getuid()` and print the result.

```
gcc test_getuid.c -o test_getuid  
./test_getuid
```


If the program returns the correct user ID, the system call is working as expected.

7. Inserting Kernel Modules (If Required)

If your implementation involves adding new kernel modules (e.g., for device drivers or other kernel extensions), you can insert them after compiling them by using the `insmod` command.

```
sudo insmod your_module.ko
```

To verify that the module is inserted, use:

```
lsmod | grep your_module
```

This command will list all loaded modules and check for the presence of your module.

Conclusion: Compiling and inserting the module is a crucial part of kernel development. After compiling the kernel and inserting any required modules, the `getuid()` system call can be tested to ensure it functions as expected. If there are no errors and the program produces the expected output (the real user ID), the implementation is complete.

11.5 Test Program Code (for `getuid()` System Call)

After successfully implementing and compiling the `getuid()` system call into the Linux kernel, you need a **user-space test program** to verify that the system call works correctly.

Below is a **simple C program** that directly invokes your newly added system call using the `syscall()` function:

Test Program: `test_getuid.c`

```
#include <stdio.h>
#include <unistd.h>
#include <sys/syscall.h>

#define __NR_getuid 233 // Replace with the syscall number you assigned in syscall_64.tbl

int main() {
    long uid;

    // Invoke the system call using syscall()
    uid = syscall(__NR_getuid);

    // Display the result
    printf("User ID returned by custom getuid syscall: %ld\n", uid);

    return 0;
}
```

How to Compile and Run:

Compile the Program:

```
gcc test_getuid.c -o test_getuid
```

Run the Program:

```
./test_getuid
```

Expected Output:

If everything is correctly implemented, the program should print your real user ID (usually 0 for root, or a non-zero UID for regular users):

User ID returned by custom getuid syscall: 1000

12:Issues and Challenges Faced During System Call Implementation

12.1: Kernel compilation failed due to missing dependencies.

12.2:Syscall number conflict.

13: Solutions

- Installed necessary build dependencies.
- Chose an unused syscall number from syscall table.

14:Evaluation Criteria Reflections

14.1 Installation Clarity (Evaluation Criteria Reflections)

The installation process of BlackArch Linux on VMware Workstation was documented with clear and detailed steps, ensuring it is understandable for both beginners and intermediate users. Each phase of the installation—from downloading the ISO to completing the setup—was broken down logically and supported with screenshots to enhance clarity.

Key Highlights:

Pre-Installation Setup: Explained how to configure VMware (CPU, memory, disk, etc.) and select the ISO image for booting.

Step-by-Step Instructions: Each installation step was presented in sequential order, making it easy to follow.

User Account Setup: Full name and account creation process were clearly shown, matching the system's final configuration.

Visual Guidance: Screenshots provided a visual reference , reducing confusion during actual installation.

Reflection:

The clear documentation helped in identifying and resolving issues during the process (e.g., dependency errors, network issues). This not only improved technical understanding but also emphasized the importance of thorough documentation in system programming and OS-level projects.

14.2 Historical Understanding (Evaluation Criteria Reflections)

The assignment demonstrated a solid understanding of the historical background of **BlackArch Linux**, tracing its roots back to **Arch Linux** and analyzing how it evolved into a powerful security-focused distribution.

Key Highlights:

Evolution from Arch Linux: Clearly explained how BlackArch builds upon Arch Linux's lightweight and flexible architecture while adding over 2,800 penetration testing tools tailored for cybersecurity professionals.

Purpose and Development Timeline: Covered the motivation behind BlackArch's development — to provide a complete suite of tools for ethical hacking, security testing, and digital forensics, with references to its release milestones.

Comparison with Other Distros: Compared BlackArch with Kali Linux and Parrot OS, emphasizing differences in package management (Pacman vs. APT), user experience, tool coverage, and flexibility.

Reflection:

Researching and writing this section deepened my understanding of how operating systems can evolve to serve specific communities (e.g., security researchers). It also illustrated the power of open-source collaboration in creating robust and specialized systems.

This historical knowledge is essential when selecting the right OS for system programming or cybersecurity projects, ensuring the tools and platform align with project goals.

14.3 Virtualization Concepts (Evaluation Criteria Reflections)

The assignment effectively explored the fundamental concepts of **virtualization** and its relevance in the context of operating systems and system-level experimentation.

Key Highlights:

Definition and Purpose: Explained virtualization as the creation of virtual instances of computing environments, allowing multiple OSes to run on a single physical machine.

How It Works: Clarified how a **hypervisor** abstracts hardware and allocates resources to virtual machines (VMs), enabling isolated and efficient OS execution.

Types of Hypervisors: Discussed both **Type 1 (bare-metal)** and **Type 2 (hosted)** hypervisors, with VMware Workstation classified as Type 2—ideal for developers and learners.

Benefits in OS Development: Emphasized how virtualization makes it safer and more practical to test OS features (like system calls) without risking the host system. Snapshots and rollbacks add further reliability.

Reflection:

Understanding virtualization was crucial for this assignment. It enabled me to safely install and experiment with BlackArch Linux, modify kernel code (e.g., for `getuid()`), and troubleshoot issues in an isolated environment. This section helped bridge theoretical OS concepts with real-world development and testing practices.

The experience has also highlighted how virtualization tools accelerate learning, reduce hardware constraints, and enable parallel experimentation—an invaluable asset for any system programmer or OS enthusiast.

14.4 System Programming Experience (Evaluation Criteria Reflections)

This assignment provided valuable hands-on experience in **system programming**, particularly through the implementation of the `getuid()` system call in the Linux kernel.

Key Highlights:

Kernel Source Editing: Gained real-world exposure to modifying kernel-level code by adding a custom system call, understanding the structure of kernel directories, and navigating low-level source files.

Syscall Table Registration: Learned how system calls are identified and mapped using syscall tables, and how system call numbers are assigned and referenced in both kernel and user space.

Compiling the Kernel: Experienced the full cycle of building the Linux kernel, including resolving dependency errors and using `make`, `make modules_install`, and `make install`.

Testing with User Programs: Wrote a C test program to call the new `getuid()` syscall using `syscall(__NR_getuid)`, confirming successful integration between kernel space and user space.

Reflection:

This part of the assignment significantly improved my confidence and capability in system-level programming. Editing and compiling a real OS kernel—especially a

security-focused distro like BlackArch—felt challenging but incredibly rewarding. It taught me the importance of precision, debugging, and deep understanding when working close to hardware or system-level abstractions.

14.5 Filesystem Knowledge (Evaluation Criteria Reflections)

Exploring the filesystem support in BlackArch Linux provided deeper insights into how operating systems manage data storage, access, and structure. This section helped solidify my understanding of Linux filesystems and their role in system performance and stability.

Key Highlights:

Supported Filesystems: Identified the various filesystems supported by BlackArch Linux, including **ext4**, **Btrfs**, **XFS**, **F2FS**, and others. Ext4 was found to be the default and most stable choice for installations.

Filesystem Selection: Discussed why **ext4** was chosen during installation—primarily for its balance between performance, reliability, journaling support, and widespread community adoption.

Technical Comparison: Briefly compared ext4 with Btrfs and XFS in terms of features like snapshot support, error detection, performance under heavy loads, and maintenance requirements.

Reflection:

This part of the assignment enhanced my appreciation for how critical filesystems are in an operating system. It also made me realize that choosing the right filesystem can impact system stability, performance, and ease of recovery. In the context of system programming, understanding how files are structured and accessed at a low level is essential when writing system calls, working with file descriptors, or debugging OS behavior.

By working with BlackArch's filesystem during installation and kernel modification, I gained practical experience that will benefit future OS development and troubleshooting tasks.

15. Conclusion:

15.1 Lessons Learned

Throughout the course of this OSSP individual assignment, I gained valuable theoretical knowledge and practical skills in operating systems and system-level programming. Some of the key lessons include:

Hands-on OS Installation: Installing BlackArch Linux on VMware gave me practical experience in setting up a virtualized environment, configuring hardware resources, and troubleshooting installation issues.

System Programming Fundamentals: Implementing a custom system call (`getuid()`) deepened my understanding of how user-space programs interact with the kernel, how system calls are defined, registered, and tested, and how kernel compilation and modules work.

Debugging and Problem-Solving: Facing real-world issues such as network errors and package dependency problems improved my troubleshooting skills and taught me the importance of logs, documentation, and online resources.

Filesystem and OS Internals: I learned about different filesystems supported by Linux and why ext4 is commonly used, reinforcing my understanding of data management within an OS.

Virtualization and Hypervisors: I developed a solid foundation in virtualization technologies and their importance in OS development, testing, and safe experimentation.

15.2 Summary of Work Done

Documented the historical background and evolution of BlackArch Linux, its relation to Arch Linux, and comparison with similar distributions like Kali Linux.

Installed BlackArch Linux on VMware Workstation, configured user accounts, and provided a step-by-step guide with screenshots.

Identified and resolved installation issues, including network and dependency errors.

Explored filesystem support, specifically the advantages of using ext4, and discussed other alternatives.

Studied virtualization concepts and explained how virtualization supports OS development and experimentation.

Implemented the `getuid()` system call, modified kernel source files, registered the syscall, compiled the kernel, and tested the system call with a custom user-space program.

Reflected on evaluation criteria, highlighting lessons learned across system programming, installation clarity, OS concepts, and algorithm analysis.

This assignment has been a crucial step in my journey as a system-level developer, combining theory with practical experience, and preparing me for more advanced OS-related projects in the future.

16. Future Outlook and Recommendations

16.1 Suggestions for New Learners

For students and beginners diving into **Operating Systems and System Programming**, here are some suggestions based on my experience:

Start with Virtualization: Begin experimenting in a virtual machine using tools like VMware or VirtualBox. This provides a safe environment for testing OS installations and kernel modifications without harming your main system.

Document Every Step: Keep a log of what you do—especially while working with installations or kernel modifications. This helps in understanding errors and revisiting solutions later.

Understand Before You Modify: Never blindly copy-paste kernel code. Try to understand what each file and function does before making changes. Reading kernel documentation or forums like Stack Overflow, Arch Wiki, or Linux Kernel Mailing Lists can help a lot.

Break and Learn: Don't be afraid to make mistakes. Some of the most valuable lessons come from debugging and fixing what went wrong.

Bridge Theory and Practice: As you learn OS concepts in class (like system calls, memory management, scheduling), try to implement or observe them in practice inside a Linux distribution.

16.2 Plans for Further System Exploration

Moving forward, I plan to explore deeper areas of system-level development and operating systems:

Kernel Module Development: Beyond system calls, I want to write kernel modules to interact with devices, explore custom drivers, or simulate system-level behaviors.

Linux Security Internals: Since BlackArch is security-focused, I'll dive into Linux security modules (LSM), AppArmor, and SELinux to understand how permissions and access control work in the kernel.

Process and Memory Management: I plan to study how Linux manages processes, memory, and I/O, and try implementing custom scheduling algorithms or memory tracking tools.

File System Implementation: As a long-term goal, I'd love to experiment with implementing a simple virtual filesystem to understand how Linux handles data storage at the lowest level.

Contributing to Open Source: Eventually, I aim to contribute to open-source OS projects, including the Linux kernel or Arch-based distros like BlackArch.

